SEAGRAMMAR — in silico pilot on the final 20 + 6 compound panel

Rule-learning framework (WP0) run end-to-end on the finalised panel: 20 training compounds, 6 sealed compounds, 42 strains, 3 replicates.

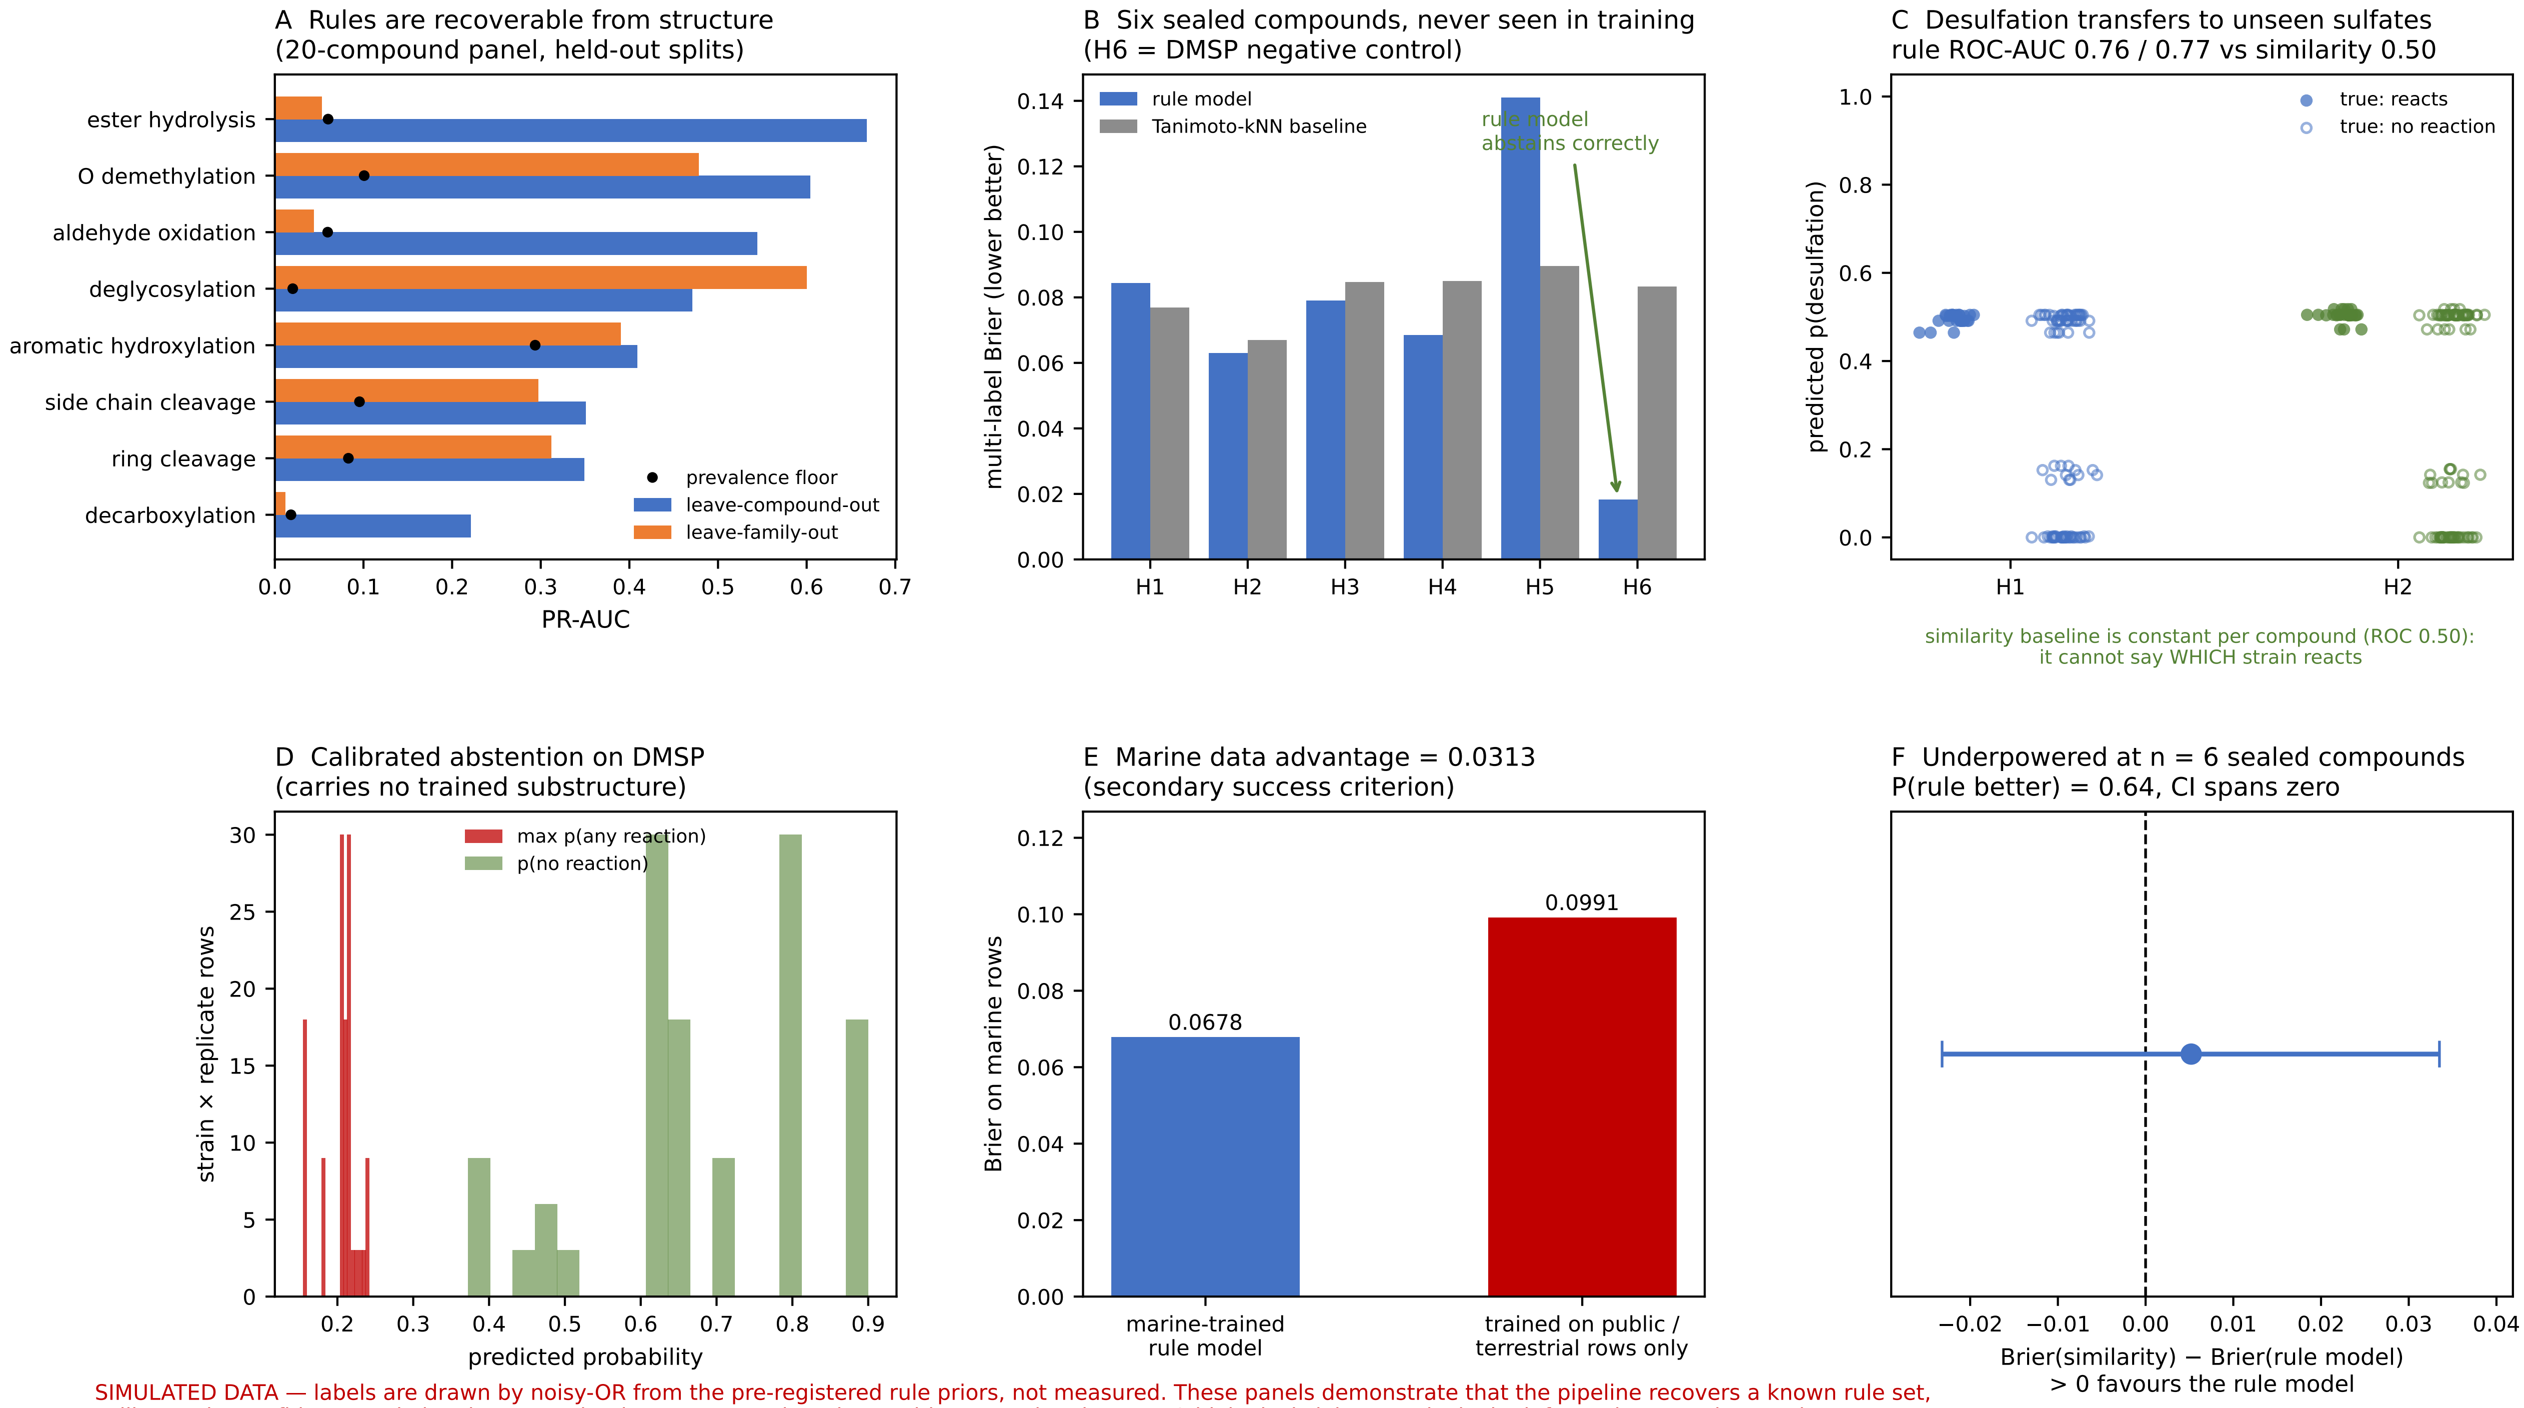

SIMULATED DATA — labels are drawn by noisy-OR from the pre-registered rule priors, not measured. These panels demonstrate that the pipeline recovers a known rule set, calibrates its confidence and abstains correctly; they are an engineering positive control and carry NO biological claim. Panel F is the informative negative result.